

Beginner's Guide to LVM (Logical Volume Management)

 thegeekdiary.com/redhat-centos-a-beginners-guide-to-lvm-logical-volume-manager/

admin

What is LVM


Logical volume manager (LVM) introduces an extra layer between the physical disks and the file system allowing file systems to be :

- resized and moved easily and online without requiring a system-wide outage.
- Using discontinuous space on disk
- meaningful names to volumes, rather than the usual cryptic device names.
- span multiple physical disks

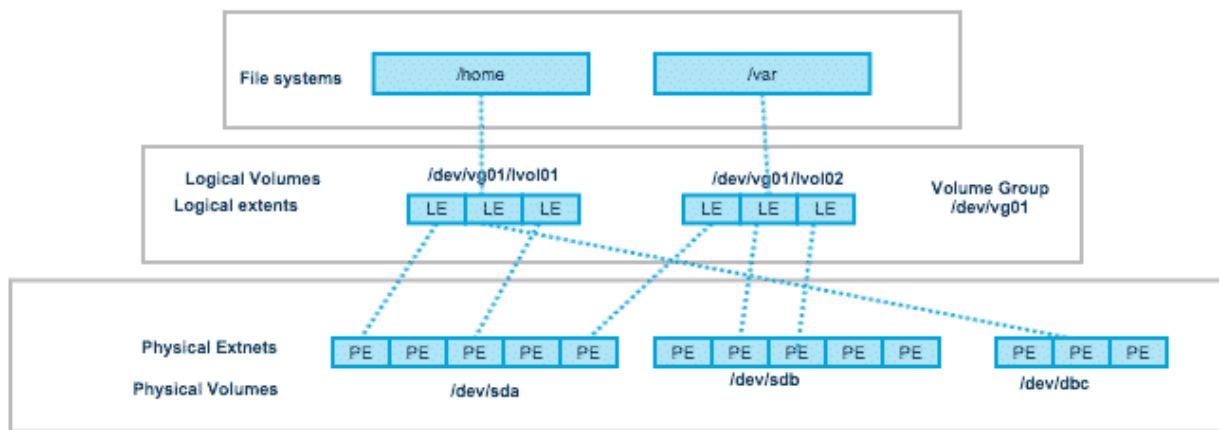
Linux LVM is very similar to HP-UX LVM and provides many other advanced features like snapshots, cluster support (GFS2, OCFS and Lustre).

Concepts

LVM comprises of few conceptual layers such as physical volume, logical volume and file systems.



The conceptual layers are in turn made up of smaller units like Physical extents(in case of Physical volumes) and Logical extents (in case of Logical Volumes).



Physical Volume (PV)

Each Physical Volume can be a disk partition, whole disk, meta-device, or a loopback file. Use the command **pvcreate** to initialize storage for use by LVM. Initializing a block device as physical volume places a label at the start of the device.

Volume Group (VG)

A Volume Group gathers together a collection of Logical Volumes and Physical Volumes into one administrative unit. Volume group is divided into fixed size physical extents. The command **vgcreate** creates a new volume group using the block special device Physical Volume path previously configured for LVM with **pvcreate**.

- VGs are made up of PVs, which in turn are made up of physical extents (PEs). The size of PE can differ in different VGs and is defined at the time of creating VG.
- The default size of PE is 4MB, but you can change it to the value you want at the time of VG creation.
- Generally, larger the PE size, better the performance (though less granular control of LV).

Logical Volume (LV)

A Logical Volume is the conceptual equivalent of a disk partition in a non-LVM system. Logical volumes are block devices which are created from the physical extents present in the same volume group. You can use command **lvcreate** to create a logical volume in an existing volume group.

File system

File systems are built on top of logical volumes. The command **mkfs** can be used to create file system on top of a logical volume. Once the file system is created we can mount the logical volume as per our need.

Lets Get Started

The example

In the example below we would :

1. Create 3 Physical volumes from 3 physical disks (**/dev/sdb**, **/dev/sdc**, **/dev/sdd**).
2. Create Volume group from these 3 PVs (**/dev/vg01**).
3. Create a Logical Volume in this VG (**/dev/vg01/lvol01**).
4. Create a File system on this LV and mount it (**/data01**).

Create Physical Volumes

The **pvcreate** command is used to initialize the PV for use by LVM. Before creating the PV, make sure the disk is visible in the OS. To scan the block devices to be used as PVs, use the **lvmdiskscan** command.

```
# lvmdiskscan
.....
/dev/sdb   [      2.00 GiB]
/dev/sdc   [      2.00 GiB]
/dev/sdd   [      2.00 GiB]
3 disks
19 partitions
0 LVM physical volume whole disks
0 LVM physical volumes
```

Initialize the block devices :

```
# pvcreate /dev/sdb /dev/sdc /dev/sdd
Physical volume "/dev/sdb" successfully created
Physical volume "/dev/sdc" successfully created
Physical volume "/dev/sdd" successfully created
```

Display physical volumes

Use the commands **pvdiskdisplay**, **pvs** and **pvscan** to display the PVs we just created.

```
# pvdisplay
"/dev/sdb" is a new physical volume of "2.00 GiB"
--- NEW Physical volume ---
PV Name           /dev/sdb
VG Name
PV Size           2.00 GiB
Allocatable       NO
PE Size           0
Total PE          0
Free PE           0
Allocated PE      0
PV UUID           Mt3F7z-a2AV-28Vn-uXe2-QejE-Z6tP-UMlQGM

"/dev/sdc" is a new physical volume of "2.00 GiB"
--- NEW Physical volume ---
PV Name           /dev/sdc
VG Name
PV Size           2.00 GiB
Allocatable       NO
PE Size           0
Total PE          0
Free PE           0
Allocated PE      0
PV UUID           5m1Fuc-yTRn-I2vG-bMfU-6SE7-53EA-s8VQjt

"/dev/sdd" is a new physical volume of "2.00 GiB"
--- NEW Physical volume ---
PV Name           /dev/sdd
VG Name
PV Size           2.00 GiB
Allocatable       NO
PE Size           0
Total PE          0
Free PE           0
Allocated PE      0
PV UUID           1x3e2A-C0Lt-DrUA-tPSM-lsMu-sn70-qg1j8p

# pvscan
PV /dev/sdb          lvm2 [2.00 GiB]
PV /dev/sdc          lvm2 [2.00 GiB]
PV /dev/sdd          lvm2 [2.00 GiB]
Total: 3 [6.00 GiB] / in use: 0 [0   ] / in no VG: 3 [6.00 GiB]

# pvs
PV          VG      Fmt  Attr PSize PFree
/dev/sdb    lvm2  a--  2.00g 2.00g
/dev/sdc    lvm2  a--  2.00g 2.00g
/dev/sdd    lvm2  a--  2.00g 2.00g
```

Create a Volume Group

Use the **vgcreate** command to create the new Volume Group **vg01** using the 3 PVs we just created. We can specify the extents with **-s** option and maximum number of PVs and LVs in the VG by using the options **-p** and **-l** respectively. All these options are optional and need not be necessarily used.

```
# vgcreate vg01 /dev/sdb /dev/sdc /dev/sdd
Volume group "vg01" successfully created
```

The optional options that are used with vgcreate command are :

Option	Meaning
-s	Physical extent size
-p	Max number of PVs
-l	Max number of LVs
-alloc	allocation policy (either contiguous, anywhere, or cling)

Displaying the VG information

The commands vgs and vgdisplay can be used to display the information about the VG we just created :

```
# vgs vg01
VG    #PV #LV #SN Attr   VSize VFree
vg01   3   0   0 wz--n- 5.99g 5.99g

# vgdisplay vg01
--- Volume group ---
VG Name                vg01
System ID
Format                 lvm2
Metadata Areas         3
Metadata Sequence No   1
VG Access               read/write
VG Status               resizable
MAX LV                 0
Cur LV                 0
Open LV                 0
Max PV                  0
Cur PV                 3
Act PV                  3
VG Size                 5.99 GiB
PE Size                 4.00 MiB
Total PE                1533
Alloc PE / Size         0 / 0
Free PE / Size          1533 / 5.99 GiB
VG UUID                 Cw7GGz-NH3o-Sax2-5jPv-buZS-938T-tmNKFa
```

Activating and deactivating VGs

The **vgchange** command can be used to activate/deactivate a volume group.

To deactivate a VG :

```
# vgchange -a n vg01
0 logical volume(s) in volume group "vg01" now active
```

To activate a VG :

```
# vgchange -a y vg01
1 logical volume(s) in volume group "vg01" now active
```

Create Logical Volume

The Logical volume can now be created in the VG using the **lvcreate** command.

- If you do not specify the LV name in the command, by default the LV is given the name **lvol#**.
- Normally if you do not specify which PV to span the LV, Logical volume will be created on the PV on a next-free basis.
- To create a logical volume lvol01 of size 5 GB :

```
# lvcreate -L 5G -n lvol01 vg01
Logical volume "lvol01" created
```

Creating a striped volume

To create a striped volume spanning all the 3 PVs we created :

```
# lvcreate -L 5G -I 4096 -i 3 -n lvol01 vg01
Rounding size (1280 extents) up to stripe boundary size (1281 extents)
Logical volume "lvol01" created
```

I - PVs to span while creating striped volume
i - stripe unit

Creating mirrored volume

To create a 3 way mirrored volume spanning the 3 PVs (sdb, sdc, sdd) :

```
# lvcreate -L 1G -m 2 -n lvol01 vg01
Logical volume "lvol01" created
```

We can also specify which devices to be used while creating the mirrored LV. In our case as we had only 3 PVs in the VG, the LV gets created by default on these 3 PVs.

Displaying the LV information

The commands **lvdisplay**, **lvs** and **lvscan** can be used to display the information about the LV we just created.

```
# lvs /dev/vg01/lvol01
LV      VG      Attr      LSize Pool Origin Data%  Move Log          Cpy%Sync Convert
lvol01  vg01    mwi-a-m-- 1.00g                lvol01_mlog    100.00
```

```
# lvdisplay /dev/vg01/lvol01
--- Logical volume ---
LV Path                /dev/vg01/lvol01
LV Name                 lvol01
VG Name                 vg01
LV UUID                 ptlmAV-m042-fwIJ-e2Ml-r9kj-PFcC-M0exxw
LV Write Access         read/write
LV Creation host, time localhost.localdomain, 2014-10-22 09:04:25 -0700
LV Status                available
# open                  0
LV Size                 1.00 GiB
Current LE              256
Mirrored volumes        3
Segments                1
Allocation              inherit
Read ahead sectors      auto
- currently set to      256
Block device            253:4

# lvscan
ACTIVE                  '/dev/vg01/lvol01' [1.00 GiB] inherit
```

Creating File system

The final step is to create a file system on the new LV we just created and mount it on a directory to be able to access it and store data in it. The command **mkfs** can be used to create file system on top of the LV.

```
# mkfs.ext4 /dev/vg01/lvol01
mke2fs 1.41.12 (17-May-2010)
Filesystem label=
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
Stride=0 blocks, Stripe width=0 blocks
65536 inodes, 262144 blocks
13107 blocks (5.00%) reserved for the super user
First data block=0
Maximum filesystem blocks=268435456
8 block groups
32768 blocks per group, 32768 fragments per group
8192 inodes per group
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376

Writing inode tables: done
Creating journal (8192 blocks): done
Writing superblocks and filesystem accounting information: done
```

This filesystem will be automatically checked every 37 mounts or 180 days, whichever comes first. Use tune2fs -c or -i to override.

The logical volume can be mounted, once the file system is created. Make sure to add an entry to **/etc/fstab**, so that it is mounted automatically when the system boots.

```
# mkdir /data01
# mount /dev/vg01/lvol01 /data01

# vi /etc/fstab
/dev/vg01/lvol01          /data01          ext4      defaults    0 0

# df -h /data01
Filesystem                Size      Used    Avail    Use%    Mounted on
/dev/mapper/vg01-lvol01   1008M     34M    924M     4%      /data01
```

The Graphical Tool to manage LVM

There is a cool graphical tool available (**system-config-lvm**) in case you want to use. If not already installed on the system, install it using yum:

```
# yum install system-config-lvm
```

To start the Graphical LVM administration tool, fire the command :

```
# system-config-lvm
```

